

# Laburen.com – Product Engineer Challenge

## Objective

Build a functional AI Agent that demonstrates your skills as an AI Product Engineer and DevOps professional.

The agent must:

- Enable knowledge management through RAG
- Persist all conversation data in PostgreSQL using Prisma
- Deliver a real-time streaming experience with visual tool states

We want to see how you think end-to-end: from code architecture → automated deployment → user experience → system observability.

## Functional Requirements

### 1. Complete Streaming Chat Experience

- Real-time conversation with streaming
- Custom streaming components showing different tool states:
  - "Searching documents..."
  - "Uploading file to RAG..."
  - "Processing embedding..."
  - Visual progress for each operation
- Clear indicators when the agent is:
  - Thinking
  - Using a tool
  - Responding
- Display model reasoning in real-time
  - Show the model's reasoning / thinking process as it streams
  - Users must see how the model thinks and reasons live
- MANDATORY: Use Vercel AI SDK with streaming

### 2. Complete Persistence with Prisma + PostgreSQL

- Persist everything:
  - User messages
  - Agent responses
  - Tool calls
  - Tool states
  - Reasoning traces

- On page reload or returning to the chat → state must be identical:
  - Same message history
  - Same tool states displayed
  - Same conversation context
  - Reasoning / thinking process preserved
- Users must be able to create NEW chats from scratch
  - Clear UI to start a new conversation
  - Ability to switch between different chats
  - Each chat maintains its own independent history and context
- Well-designed Prisma schema covering:
  - Conversations / Chats
  - Messages (role, content, tool\_calls, reasoning, etc.)
  - Tool states
  - Streaming metadata

### **3. RAG Tools (2 required)**

#### **Tool 1: Upload to RAG**

- Upload files (PDF, TXT, MD, DOCX) to the RAG system
- Process and generate embeddings
- Save to vector database
- Show progress in streaming UI with custom component

#### **Tool 2: Search in RAG**

- Search for information in indexed documents
- Display sources / references clearly:
  - Document name
  - Exact fragment / chunk
  - Similarity score
  - Chunk metadata (if available)
- Render sources in a clean, readable way in the chat

Note on RAG: Use any vector database you prefer (pgvector, Pinecone, Chroma, Qdrant, etc.). What matters is that it works reliably.

### **4. User Experience**

- Clear visualization of:
  - User messages
  - Agent responses
  - Tool calls in progress (custom components)
  - Information sources
  - Model reasoning / thinking process
- Well-handled loading states and errors
- Ability to create and switch between multiple chats
- Responsive and accessible design



## Required Tech Stack

| Category        | Required Technology                         |
|-----------------|---|
| Frontend        | Next.js (Pages Router)                      |
| AI Framework    | Vercel AI SDK (mandatory)                   |
| UI Components   | Vercel AI Elements (mandatory)              |
| Database        | PostgreSQL                                  |
| ORM             | Prisma                                      |
| Vector Database | Your choice (pgvector, Pinecone, Chroma...) |
| LLM Provider    | OpenAI, Anthropic, Azure OpenAI, Ollama...  |



## DevOps Requirements

### Task 1: Deploy with Maximum Automation

Show everything you know about DevOps in the deployment process.

Document in the README:

- Which automations you implemented
- Why you chose each tool / service
- How the deployment flow works

### Task 2: DevOps Reflection on the Agent

After completing the challenge, write a document (`devops-proposal.md`) where you reflect and propose improvements for production.

Guiding Questions:

1. Observability: How would you monitor this agent in production? (metrics, logs, traces)
2. Scalability: What would you do if traffic increases 1000×?
3. Reliability: How would you ensure 99.9% uptime?
4. Cost optimization: How would you optimize costs for LLM calls, vector DB, and hosting?
5. Security: What vulnerabilities do you see and how would you mitigate them?
6. Testing: What testing strategy would you propose for this kind of system?
7. Advanced CI/CD: What improvements would you make to the pipeline you built?
8. Disaster recovery: What backup and recovery plan would you implement?

Expected format:

- Analysis of the current architecture
- Concrete proposals with pros/cons
- Diagrams if helpful (architecture, flows, etc.)
- Proposed tech stack for production
- Prioritized implementation roadmap

This is not mandatory to implement — we just want to see how you think as a DevOps engineer when facing an AI system in production.

## Evaluation Criteria

| Criterion                      | Weight | What we evaluate   |
|--------------------------------|--------|--|
| Architecture & Code Quality    | 15%    | Clarity, patterns, organization, Prisma schema                       |
| AI SDK + Tools Integration     | 15%    | Streaming, tool calling, state management, reasoning display         |
| RAG + Source Citation          | 10%    | Effective retrieval, clear & useful source presentation              |
| Persistence & State Management | 10%    | Prisma usage, conversation recovery, consistency, multi-chat support |
| UX / Streaming Experience      | 10%    | Custom components, fluid feel, reasoning visualization               |
| DevOps & Deployment            | 15%    | Automation, CI/CD, functional public deployment                      |
| DevOps Reflection              | 10%    | Quality of analysis and realistic proposals                          |
| Conceptual Design Document     | 10%    | Quality and clarity of conceptual_design.md                          |
| Extras / Creativity            | 5%     | Bonus points for thoughtful additions                                |

## Estimated Time

7 calendar days from receiving the assignment.

Focus on delivering a functional MVP that covers the core requirements. Extras are welcome but optional.



## Final Delivery

### 1. GitHub Repository

Must include:

- Complete, well-documented code
- README containing:
  - Setup instructions
  - Tech stack overview
  - Key architecture decisions
  - How to run locally
  - How to deploy
- `devops-proposal.md` — your DevOps reflection
- `conceptual_design.md` —(MANDATORY):

It must describe the architecture of the agent you are building.

It must include:

- 1 Explanation of the design approach
- 2 Reasoning and execution flow of the agent
- 3 Definition of tools
- 4 Diagrams using **Mermaid** (Markdown-based diagram syntax) representing:
  - System architecture
  - Agent flow

- `.env.example` file

### 2. Demo Video (3–5 minutes)

Record (Loom, YouTube, etc.) showing:

- Full agent usage flow
- Uploading a document to RAG
- Searching with visible sources
- Page reload → preserved conversation
- Creating a new chat & switching between chats
- Streaming components & tool states in action
- Model reasoning / thinking process visible live
- Deployment pipeline (optional but adds value)

### 3. Public Deploy

- Working public URL to test the agent
- Must be live and accessible
- HTTPS enabled (free on Vercel, Railway, Render, etc.)



## What We're Looking For in an AI Product Engineer

1. Product vision — you understand the problem and build a coherent solution
2. AI-first integration — comfortable with LLMs, streaming, tools & RAG
3. Full-stack ownership — frontend + backend + database + AI working together
4. DevOps mindset — thinking about deployment, automation, observability
5. Technical judgment — justified decisions, not just following tutorials
6. Execution speed — delivering a solid MVP under time pressure

Good luck!

We're excited to see what you build. 